

# Student Grades Prediction

## Project Description

The dataset contains grades scored by students throughout their university tenure in various courses and their CGPA calculated based on their grades Columns Description- total 43 columns -Seat No : The enrolled number of candidate that took the exams

CGPA : The cumulative GPA based on the four year total grade progress of each candidate .  
CGPA is a Final Marks -- provided to student.

· All other columns are course codes in the format AB-XXX where AB are alphabets representing candidates' departments and XXX are numbers where first X represents the year the candidate took exam

**Predict** - CGPA of a student based on different grades in four years.

# Importing Libraries

```
In [1]: # Analyse and Manipulate Data
import numpy as np
import pandas as pd

# Data Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Preprocessor
from sklearn.preprocessing import StandardScaler

#Model Evaluation
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV

# Regressor
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor,
from xgboost import XGBRegressor

#Metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Model saving
import pickle

# Prevent WARNINGS!
import warnings
warnings.filterwarnings("ignore")
```

## Import & Analyse Data

```
In [2]: df=pd.read_csv("E:\\dataset\\dataset4-main\\Grades.csv")
df
```

Out[2]:

|     | Seat<br>No.  | PH-<br>121 | HS-<br>101 | CY-<br>105 | HS-<br>105/12 | MT-<br>111 | CS-<br>105 | CS-<br>106 | EL-<br>102 | EE-<br>119 | ... | CS-<br>312 | CS-<br>317 | CS-<br>403 | CS-<br>421 | CS-<br>406 | CS-<br>414 |
|-----|--------------|------------|------------|------------|---------------|------------|------------|------------|------------|------------|-----|------------|------------|------------|------------|------------|------------|
| 0   | CS-<br>97001 | B-         | D+         | C-         | C             | C-         | D+         | D          | C-         | B-         | ... | C-         | C-         | C-         | C-         | A-         | A          |
| 1   | CS-<br>97002 | A          | D          | D+         | D             | B-         | C          | D          | A          | D+         | ... | D+         | D          | C          | D          | A-         | B-         |
| 2   | CS-<br>97003 | A          | B          | A          | B-            | B+         | A          | B-         | B+         | A-         | ... | B          | B          | A          | C          | A          | A          |
| 3   | CS-<br>97004 | D          | C+         | D+         | D             | D          | A-         | D+         | C-         | D          | ... | D+         | C          | D+         | C-         | B-         | B          |
| 4   | CS-<br>97005 | A-         | A-         | A-         | B+            | A          | A          | A-         | B+         | A          | ... | B-         | B+         | B+         | B-         | A-         | A          |
| ... | ...          | ...        | ...        | ...        | ...           | ...        | ...        | ...        | ...        | ...        | ... | ...        | ...        | ...        | ...        | ...        | ...        |
| 566 | CS-<br>97567 | B          | A          | A          | A-            | A+         | A          | A-         | A-         | A+         | ... | A-         | A-         | A          | A          | A          | B+         |
| 567 | CS-<br>97568 | A+         | A          | A          | A             | A          | A          | A          | A-         | A          | ... | B+         | B+         | A          | A          | A-         | B          |
| 568 | CS-<br>97569 | B          | A          | A-         | B+            | A          | A          | A          | A          | A          | ... | A-         | B          | A          | B+         | A          | C          |
| 569 | CS-<br>97570 | A          | B+         | D          | A             | D          | D+         | B-         | C-         | B-         | ... | D          | B          | B          | C-         | D          | C          |
| 570 | CS-<br>97571 | C          | D          | D          | C             | C          | D+         | B          | C+         | C          | ... | C+         | C          | B-         | D          | F          | C-         |

571 rows × 43 columns




# Data Inspection

```
In [3]: df.head()
```

Out[3]:

|   | Seat No. | PH-121 | HS-101 | CY-105 | HS-105/12 | MT-111 | CS-105 | CS-106 | EL-102 | EE-119 | ... | CS-312 | CS-317 | CS-403 | CS-421 | CS-406 | CS-414 | CS-414 |
|---|----------|--------|--------|--------|-----------|--------|--------|--------|--------|--------|-----|--------|--------|--------|--------|--------|--------|--------|
| 0 | CS-97001 | B-     | D+     | C-     | C         | C-     | D+     | D      | C-     | B-     | ... | C-     | C-     | C-     | C-     | A-     | A      | C      |
| 1 | CS-97002 | A      | D      | D+     | D         | B-     | C      | D      | A      | D+     | ... | D+     | D      | C      | D      | A-     | B-     | C      |
| 2 | CS-97003 | A      | B      | A      | B-        | B+     | A      | B-     | B+     | A-     | ... | B      | B      | A      | C      | A      | A      | C      |
| 3 | CS-97004 | D      | C+     | D+     | D         | D      | A-     | D+     | C-     | D      | ... | D+     | C      | D+     | C-     | B-     | B      | C      |
| 4 | CS-97005 | A-     | A-     | A-     | B+        | A      | A      | A-     | B+     | A      | ... | B-     | B+     | B+     | B-     | A-     | A      | C      |

5 rows × 43 columns




```
In [4]: df.tail()
```

Out[4]:

|     | Seat No. | PH-121 | HS-101 | CY-105 | HS-105/12 | MT-111 | CS-105 | CS-106 | EL-102 | EE-119 | ... | CS-312 | CS-317 | CS-403 | CS-421 | CS-406 | CS-414 | CS-414 |
|-----|----------|--------|--------|--------|-----------|--------|--------|--------|--------|--------|-----|--------|--------|--------|--------|--------|--------|--------|
| 566 | CS-97567 | B      | A      | A      | A-        | A+     | A      | A-     | A-     | A+     | ... | A-     | A-     | A      | A      | A      | A      | B+     |
| 567 | CS-97568 | A+     | A      | A      | A         | A      | A      | A      | A-     | A      | ... | B+     | B+     | A      | A      | A-     | B      | C      |
| 568 | CS-97569 | B      | A      | A-     | B+        | A      | A      | A      | A      | A      | ... | A-     | B      | A      | B+     | A      | C      | C      |
| 569 | CS-97570 | A      | B+     | D      | A         | D      | D+     | B-     | C-     | B-     | ... | D      | B      | B      | C-     | D      | C      | C      |
| 570 | CS-97571 | C      | D      | D      | C         | C      | D+     | B      | C+     | C      | ... | C+     | C      | B-     | D      | F      | C-     | C      |

5 rows × 43 columns



```
In [5]: # Dimension of data
print("Dataset contain {0} rows & {1} columns".format(df.shape[0],df.shape[1]))
Dataset contain 571 rows & 43 columns
```

```
In [6]: # Name of the columns
print("Columns/Variables we have in our dataset are:\n\n",df.columns)
```

Columns/Variables we have in our dataset are:

```
Index(['Seat No.', 'PH-121', 'HS-101', 'CY-105', 'HS-105/12', 'MT-111',
      'CS-105', 'CS-106', 'EL-102', 'EE-119', 'ME-107', 'CS-107', 'HS-205/2
0',
      'MT-222', 'EE-222', 'MT-224', 'CS-210', 'CS-211', 'CS-203', 'CS-214',
      'EE-217', 'CS-212', 'CS-215', 'MT-331', 'EF-303', 'HS-304', 'CS-301',
      'CS-302', 'TC-383', 'MT-442', 'EL-332', 'CS-318', 'CS-306', 'CS-312',
      'CS-317', 'CS-403', 'CS-421', 'CS-406', 'CS-414', 'CS-419', 'CS-423',
      'CS-412', 'CGPA'],
      dtype='object')
```

```
In [7]: #Data info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 571 entries, 0 to 570
Data columns (total 43 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Seat No.        571 non-null   object
1   PH-121          571 non-null   object
2   HS-101          571 non-null   object
3   CY-105          570 non-null   object
4   HS-105/12       570 non-null   object
5   MT-111          569 non-null   object
6   CS-105          571 non-null   object
7   CS-106          569 non-null   object
8   EL-102          569 non-null   object
9   EE-119          569 non-null   object
10  ME-107          569 non-null   object
11  CS-107          569 non-null   object
12  HS-205/20       566 non-null   object
13  MT-222          566 non-null   object
14  EE-222          564 non-null   object
15  MT-224          564 non-null   object
16  CS-210          564 non-null   object
17  CS-211          566 non-null   object
18  CS-203          566 non-null   object
19  CS-214          565 non-null   object
20  EE-217          565 non-null   object
21  CS-212          565 non-null   object
22  CS-215          565 non-null   object
23  MT-331          562 non-null   object
24  EF-303          561 non-null   object
25  HS-304          561 non-null   object
26  CS-301          561 non-null   object
27  CS-302          561 non-null   object
28  TC-383          561 non-null   object
29  MT-442          561 non-null   object
30  EL-332          562 non-null   object
31  CS-318          562 non-null   object
32  CS-306          562 non-null   object
33  CS-312          561 non-null   object
34  CS-317          559 non-null   object
35  CS-403          559 non-null   object
36  CS-421          559 non-null   object
37  CS-406          486 non-null   object
38  CS-414          558 non-null   object
39  CS-419          558 non-null   object
40  CS-423          557 non-null   object
41  CS-412          492 non-null   object
42  CGPA            571 non-null   float64
dtypes: float64(1), object(42)
memory usage: 191.9+ KB
```

```
In [8]: # Datatypes of the columns
df.dtypes
```

```
Out[8]: Seat No.      object
PH-121      object
HS-101      object
CY-105      object
HS-105/12   object
MT-111      object
CS-105      object
CS-106      object
EL-102      object
EE-119      object
ME-107      object
CS-107      object
HS-205/20   object
MT-222      object
EE-222      object
MT-224      object
CS-210      object
CS-211      object
CS-203      object
CS-214      object
EE-217      object
CS-212      object
CS-215      object
MT-331      object
EF-303      object
HS-304      object
CS-301      object
CS-302      object
TC-383      object
MT-442      object
EL-332      object
CS-318      object
CS-306      object
CS-312      object
CS-317      object
CS-403      object
CS-421      object
CS-406      object
CS-414      object
CS-419      object
CS-423      object
CS-412      object
CGPA        float64
dtype: object
```

```
In [9]: #renaming two columns for our convenience
df.rename(columns={'HS-105/12': 'HS-105', 'HS-205/20': 'HS-205'},inplace=True)
```

## Observations:

- We can clearly see that in the dataset there are 43 columns and 571 rows.

- Renamed two columns for our convenience
- We have two types of datatypes(object and float) present in our dataset.
- As we can clearly see that all the columns in our dataset has object values except CGPA column which is our target/label.
- Memory usage: 191.9+ KB

## Duplicate & Missing Values

```
In [10]: #Checking for Duplicate Values  
print("We have {} duplicated values in our dataframe".format(df.duplicated().sum()))
```

We have 0 duplicated values in our dataframe

```
In [11]: #checking for missing Values  
  
#Checking Null Values  
df.isnull().sum().sum()
```

Out[11]: 425



```
In [12]: df.isnull().sum()
```

```
Out[12]: Seat No.      0
         PH-121      0
         HS-101      0
         CY-105      1
         HS-105      1
         MT-111      2
         CS-105      0
         CS-106      2
         EL-102      2
         EE-119      2
         ME-107      2
         CS-107      2
         HS-205      5
         MT-222      5
         EE-222      7
         MT-224      7
         CS-210      7
         CS-211      5
         CS-203      5
         CS-214      6
         EE-217      6
         CS-212      6
         CS-215      6
         MT-331      9
         EF-303     10
         HS-304     10
         CS-301     10
         CS-302     10
         TC-383     10
         MT-442     10
         EL-332      9
         CS-318      9
         CS-306      9
         CS-312     10
         CS-317     12
         CS-403     12
         CS-421     12
         CS-406     85
         CS-414     13
         CS-419     13
         CS-423     14
         CS-412     79
         CGPA        0
         dtype: int64
```

## Handling Missing Values

```
In [13]: def replace_null_with_mode(df):
# Find columns with null values
columns_with_null = df.columns[df.isnull().any()]

# Iterate over columns with null values
for column in columns_with_null:
# Calculate mode
mode_value = df[column].mode()[0]

# Replace null values with mode
df[column].fillna(mode_value, inplace=True)

return df
```

```
In [14]: df= replace_null_with_mode(df)
df
```

```
Out[14]:
```

|     | Seat<br>No. | PH-<br>121 | HS-<br>101 | CY-<br>105 | HS-<br>105 | MT-<br>111 | CS-<br>105 | CS-<br>106 | EL-<br>102 | EE-<br>119 | ... | CS-<br>312 | CS-<br>317 | CS-<br>403 | CS-<br>421 | CS-<br>406 | CS-<br>414 | CS-<br>415 |
|-----|-------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|-----|------------|------------|------------|------------|------------|------------|------------|
| 0   | CS-97001    | B-         | D+         | C-         | C          | C-         | D+         | D          | C-         | B-         | ... | C-         | C-         | C-         | C-         | A-         | A          | C          |
| 1   | CS-97002    | A          | D          | D+         | D          | B-         | C          | D          | A          | D+         | ... | D+         | D          | C          | D          | A-         | B-         |            |
| 2   | CS-97003    | A          | B          | A          | B-         | B+         | A          | B-         | B+         | A-         | ... | B          | B          | A          | C          | A          | A          |            |
| 3   | CS-97004    | D          | C+         | D+         | D          | D          | A-         | D+         | C-         | D          | ... | D+         | C          | D+         | C-         | B-         | B          | C          |
| 4   | CS-97005    | A-         | A-         | A-         | B+         | A          | A          | A-         | B+         | A          | ... | B-         | B+         | B+         | B-         | A-         | A          | A          |
| ... | ...         | ...        | ...        | ...        | ...        | ...        | ...        | ...        | ...        | ...        | ... | ...        | ...        | ...        | ...        | ...        | ...        | ...        |
| 566 | CS-97567    | B          | A          | A          | A-         | A+         | A          | A-         | A-         | A+         | ... | A-         | A-         | A          | A          | A          | B+         | B          |
| 567 | CS-97568    | A+         | A          | A          | A          | A          | A          | A          | A-         | A          | ... | B+         | B+         | A          | A          | A-         | B          | A          |
| 568 | CS-97569    | B          | A          | A-         | B+         | A          | A          | A          | A          | A          | ... | A-         | B          | A          | B+         | A          | C          | B          |
| 569 | CS-97570    | A          | B+         | D          | A          | D          | D+         | B-         | C-         | B-         | ... | D          | B          | B          | C-         | D          | C          |            |
| 570 | CS-97571    | C          | D          | D          | C          | C          | D+         | B          | C+         | C          | ... | C+         | C          | B-         | D          | F          | C-         | B          |

571 rows × 43 columns




```
In [15]: #Checking Null values  
df.isnull().sum()
```

```
Out[15]: Seat No.      0  
PH-121      0  
HS-101      0  
CY-105      0  
HS-105      0  
MT-111      0  
CS-105      0  
CS-106      0  
EL-102      0  
EE-119      0  
ME-107      0  
CS-107      0  
HS-205      0  
MT-222      0  
EE-222      0  
MT-224      0  
CS-210      0  
CS-211      0  
CS-203      0  
CS-214      0  
EE-217      0  
CS-212      0  
CS-215      0  
MT-331      0  
EF-303      0  
HS-304      0  
CS-301      0  
CS-302      0  
TC-383      0  
MT-442      0  
EL-332      0  
CS-318      0  
CS-306      0  
CS-312      0  
CS-317      0  
CS-403      0  
CS-421      0  
CS-406      0  
CS-414      0  
CS-419      0  
CS-423      0  
CS-412      0  
CGPA        0  
dtype: int64
```

```
In [16]: #Visualizing null values
sns.heatmap(df.isnull())
```

Out[16]: <Axes: >



### Observations:

- **Duplicate Values** Our dataset does not contain any duplicate values.
- **Missing Values:** There are a total of 425 missing values distributed across various columns. And we have filled those missing values with mode values. So now we don't have any missing values in our dataframe.

```
In [17]: # checking value counts
for i in df.columns:
    print(df[i].value_counts())
```

CS-97001 1  
CS-97384 1  
CS-97378 1  
CS-97379 1  
CS-97380 1

..  
CS-97185 1  
CS-97184 1  
CS-97183 1  
CS-97182 1  
CS-97571 1

Name: Seat No., Length: 571, dtype: int64

A- 112  
A 111  
B+ 61  
B 57  
B- 56  
D 44  
C 33  
C+ 31  
D+ 22  
A+ 22  
C- 19  
WU 2  
F 1

Name: PH-121, dtype: int64

A- 82  
B- 78  
C 68  
B 63  
B+ 59  
C- 50  
C+ 47  
D 45  
A 38  
D+ 36  
A+ 4  
F 1

Name: HS-101, dtype: int64

A 178  
A- 120  
B+ 50  
B 49  
B- 42  
D 31  
A+ 31  
C 19  
C+ 17  
C- 16  
D+ 14  
WU 3  
F 1

Name: CY-105, dtype: int64

A 97  
A- 75  
B+ 70  
B 57

|    |    |
|----|----|
| D  | 45 |
| C  | 41 |
| B- | 40 |
| C+ | 39 |
| C- | 36 |
| D+ | 34 |
| A+ | 34 |
| WU | 2  |
| F  | 1  |

Name: HS-105, dtype: int64

|    |     |
|----|-----|
| A- | 107 |
| A  | 100 |
| B- | 70  |
| B+ | 62  |
| B  | 55  |
| C- | 39  |
| C+ | 33  |
| C  | 30  |
| D  | 26  |
| A+ | 23  |
| D+ | 21  |
| WU | 3   |
| F  | 2   |

Name: MT-111, dtype: int64

|    |     |
|----|-----|
| A  | 151 |
| A- | 134 |
| B+ | 60  |
| B  | 51  |
| A+ | 43  |
| B- | 38  |
| C+ | 23  |
| C  | 22  |
| C- | 22  |
| D+ | 15  |
| D  | 12  |

Name: CS-105, dtype: int64

|    |     |
|----|-----|
| A- | 118 |
| B+ | 101 |
| B  | 96  |
| A  | 56  |
| B- | 54  |
| C+ | 41  |
| D+ | 29  |
| C- | 27  |
| D  | 24  |
| C  | 18  |
| A+ | 4   |
| WU | 2   |
| F  | 1   |

Name: CS-106, dtype: int64

|    |     |
|----|-----|
| A- | 107 |
| A  | 92  |
| B+ | 69  |
| B  | 59  |
| B- | 53  |
| D  | 38  |
| C+ | 35  |

|                            |     |
|----------------------------|-----|
| C-                         | 32  |
| C                          | 30  |
| A+                         | 29  |
| D+                         | 23  |
| WU                         | 3   |
| F                          | 1   |
| Name: EL-102, dtype: int64 |     |
| A-                         | 139 |
| B+                         | 83  |
| B                          | 77  |
| A                          | 68  |
| B-                         | 48  |
| C                          | 48  |
| C+                         | 38  |
| D+                         | 26  |
| C-                         | 26  |
| D                          | 11  |
| A+                         | 6   |
| WU                         | 1   |
| Name: EE-119, dtype: int64 |     |
| A-                         | 81  |
| A                          | 77  |
| B+                         | 68  |
| D                          | 56  |
| B-                         | 56  |
| B                          | 50  |
| C                          | 49  |
| C-                         | 48  |
| C+                         | 37  |
| D+                         | 37  |
| A+                         | 8   |
| WU                         | 2   |
| F                          | 2   |
| Name: ME-107, dtype: int64 |     |
| A                          | 107 |
| A-                         | 81  |
| B+                         | 57  |
| B                          | 55  |
| C-                         | 49  |
| B-                         | 43  |
| A+                         | 42  |
| D                          | 38  |
| C+                         | 34  |
| D+                         | 31  |
| C                          | 30  |
| WU                         | 2   |
| I                          | 1   |
| F                          | 1   |
| Name: CS-107, dtype: int64 |     |
| A-                         | 155 |
| A                          | 118 |
| B                          | 97  |
| B+                         | 89  |
| B-                         | 36  |
| C+                         | 33  |
| C                          | 15  |
| C-                         | 11  |



|    |   |
|----|---|
| D+ | 9 |
| D  | 3 |
| A+ | 2 |
| F  | 2 |
| WU | 1 |

Name: HS-205, dtype: int64

|    |    |
|----|----|
| A- | 91 |
| A  | 80 |
| D  | 66 |
| B  | 61 |
| B- | 52 |
| D+ | 46 |
| C  | 43 |
| B+ | 42 |
| C+ | 39 |
| C- | 30 |
| A+ | 16 |
| F  | 3  |
| WU | 1  |
| W  | 1  |

Name: MT-222, dtype: int64

|    |     |
|----|-----|
| A  | 129 |
| A- | 121 |
| B+ | 65  |
| B  | 53  |
| C  | 39  |
| B- | 35  |
| A+ | 32  |
| C+ | 29  |
| D  | 25  |
| D+ | 21  |
| C- | 16  |
| F  | 4   |
| W  | 2   |

Name: EE-222, dtype: int64

|    |     |
|----|-----|
| A- | 127 |
| A  | 80  |
| B+ | 65  |
| B  | 57  |
| B- | 49  |
| C- | 43  |
| C+ | 39  |
| D+ | 37  |
| D  | 31  |
| C  | 30  |
| A+ | 10  |
| WU | 1   |
| W  | 1   |
| F  | 1   |

Name: MT-224, dtype: int64

|    |     |
|----|-----|
| A- | 140 |
| A  | 101 |
| B+ | 84  |
| B  | 59  |
| B- | 58  |
| C  | 30  |
| C+ | 27  |

|    |    |
|----|----|
| C- | 24 |
| D+ | 21 |
| D  | 12 |
| A+ | 12 |
| WU | 1  |
| W  | 1  |
| F  | 1  |

Name: CS-210, dtype: int64

|    |    |
|----|----|
| A- | 73 |
| A  | 67 |
| B- | 60 |
| D+ | 56 |
| B+ | 56 |
| B  | 56 |
| C+ | 55 |
| C- | 50 |
| C  | 39 |
| D  | 33 |
| A+ | 21 |
| F  | 3  |
| WU | 1  |
| W  | 1  |

Name: CS-211, dtype: int64

|    |    |
|----|----|
| A- | 93 |
| A  | 81 |
| B  | 66 |
| C+ | 62 |
| B+ | 59 |
| B- | 53 |
| D+ | 39 |
| C  | 35 |
| C- | 35 |
| D  | 30 |
| A+ | 15 |
| F  | 2  |
| I  | 1  |

Name: CS-203, dtype: int64

|    |    |
|----|----|
| C  | 82 |
| A- | 73 |
| B  | 63 |
| C- | 57 |
| B- | 56 |
| A  | 56 |
| D+ | 47 |
| C+ | 46 |
| B+ | 45 |
| D  | 31 |
| A+ | 12 |
| F  | 2  |
| I  | 1  |

Name: CS-214, dtype: int64

|    |     |
|----|-----|
| A- | 143 |
| A  | 97  |
| B+ | 70  |
| B  | 63  |
| B- | 57  |
| C  | 36  |

|    |    |
|----|----|
| C+ | 29 |
| A+ | 22 |
| C- | 20 |
| D+ | 19 |
| D  | 12 |
| F  | 2  |
| WU | 1  |

Name: EE-217, dtype: int64

|    |     |
|----|-----|
| A- | 107 |
| B+ | 86  |
| B  | 81  |
| B- | 65  |
| C  | 44  |
| A  | 43  |
| D+ | 36  |
| C+ | 35  |
| C- | 35  |
| D  | 33  |
| A+ | 4   |
| WU | 2   |

Name: CS-212, dtype: int64

|    |    |
|----|----|
| A- | 85 |
| A  | 68 |
| B  | 64 |
| B- | 59 |
| C+ | 50 |
| B+ | 50 |
| C  | 48 |
| C- | 47 |
| D  | 42 |
| D+ | 38 |
| A+ | 17 |
| WU | 1  |
| W  | 1  |
| F  | 1  |

Name: CS-215, dtype: int64

|    |     |
|----|-----|
| A  | 127 |
| A- | 103 |
| B+ | 64  |
| B  | 57  |
| B- | 46  |
| C- | 32  |
| D+ | 31  |
| A+ | 30  |
| C  | 28  |
| C+ | 27  |
| D  | 22  |
| F  | 4   |

Name: MT-331, dtype: int64

|    |     |
|----|-----|
| B  | 122 |
| B- | 92  |
| C  | 61  |
| B+ | 59  |
| C+ | 58  |
| C- | 56  |
| D+ | 49  |
| A- | 38  |

|    |    |
|----|----|
| D  | 19 |
| A  | 14 |
| F  | 2  |
| WU | 1  |

Name: EF-303, dtype: int64

|    |     |
|----|-----|
| A- | 138 |
| B  | 72  |
| B- | 70  |
| B+ | 66  |
| C  | 58  |
| C+ | 53  |
| C- | 33  |
| A  | 28  |
| D  | 20  |
| D+ | 19  |
| F  | 6   |
| A+ | 4   |
| WU | 2   |
| W  | 2   |

Name: HS-304, dtype: int64

|    |     |
|----|-----|
| A- | 118 |
| B+ | 74  |
| B  | 71  |
| A  | 66  |
| B- | 60  |
| C  | 41  |
| C+ | 37  |
| C- | 36  |
| D  | 29  |
| D+ | 29  |
| A+ | 9   |
| F  | 1   |

Name: CS-301, dtype: int64

|    |     |
|----|-----|
| A- | 123 |
| B  | 102 |
| A  | 86  |
| B+ | 81  |
| B- | 60  |
| C+ | 32  |
| C  | 28  |
| D  | 21  |
| C- | 19  |
| D+ | 10  |
| A+ | 9   |

Name: CS-302, dtype: int64

|    |     |
|----|-----|
| A  | 115 |
| A- | 73  |
| B+ | 68  |
| B  | 59  |
| C+ | 44  |
| D+ | 44  |
| C- | 42  |
| C  | 42  |
| B- | 40  |
| A+ | 23  |
| D  | 20  |
| F  | 1   |

Name: TC-383, dtype: int64

A- 150

A 130

B+ 65

B 47

A+ 39

B- 30

C- 28

C+ 24

C 24

D 20

D+ 13

F 1

Name: MT-442, dtype: int64

A- 105

A 76

B+ 68

B 67

B- 62

C 49

C+ 38

C- 32

D+ 22

D 22

A+ 20

F 9

WU 1

Name: EL-332, dtype: int64

A- 98

B- 69

B 68

B+ 65

C 53

C+ 49

A 42

C- 40

D 36

D+ 29

F 10

A+ 6

WU 5

W 1

Name: CS-318, dtype: int64

A- 129

B+ 75

A 74

B- 64

B 53

C- 40

C 37

C+ 36

D 31

D+ 18

A+ 10

F 3

WU 1

Name: CS-306, dtype: int64

|    |     |
|----|-----|
| A+ | 103 |
| A  | 86  |
| A- | 63  |
| D+ | 52  |
| C  | 48  |
| C- | 44  |
| B+ | 42  |
| B  | 37  |
| C+ | 37  |
| B- | 36  |
| D  | 19  |
| F  | 2   |
| W  | 1   |
| WU | 1   |

Name: CS-312, dtype: int64

|    |    |
|----|----|
| B- | 91 |
| B  | 75 |
| C  | 70 |
| A- | 66 |
| B+ | 58 |
| C+ | 57 |
| A  | 47 |
| C- | 41 |
| D+ | 34 |
| D  | 16 |
| A+ | 9  |
| F  | 7  |

Name: CS-317, dtype: int64

|    |     |
|----|-----|
| A  | 145 |
| A- | 106 |
| B+ | 62  |
| B- | 55  |
| B  | 54  |
| D+ | 35  |
| C  | 33  |
| C+ | 32  |
| C- | 30  |
| A+ | 15  |
| D  | 4   |

Name: CS-403, dtype: int64

|    |    |
|----|----|
| B  | 98 |
| B- | 74 |
| C  | 68 |
| C+ | 61 |
| C- | 60 |
| B+ | 60 |
| A- | 47 |
| D+ | 36 |
| A  | 25 |
| D  | 21 |
| F  | 17 |
| A+ | 2  |
| W  | 2  |

Name: CS-421, dtype: int64

|    |     |
|----|-----|
| A- | 262 |
| A  | 79  |
| B+ | 64  |

|    |    |
|----|----|
| B  | 58 |
| B- | 22 |
| C+ | 22 |
| C  | 19 |
| D+ | 14 |
| C- | 8  |
| A+ | 8  |
| F  | 6  |
| D  | 5  |
| W  | 3  |
| WU | 1  |

Name: CS-406, dtype: int64

|    |     |
|----|-----|
| A  | 189 |
| A- | 156 |
| B+ | 62  |
| B  | 54  |
| B- | 21  |
| A+ | 21  |
| C+ | 21  |
| C  | 20  |
| C- | 12  |
| F  | 7   |
| D+ | 3   |
| D  | 3   |
| W  | 2   |

Name: CS-414, dtype: int64

|    |     |
|----|-----|
| A- | 133 |
| B  | 89  |
| B+ | 85  |
| B- | 78  |
| A  | 56  |
| C+ | 46  |
| C  | 40  |
| C- | 20  |
| D+ | 13  |
| D  | 7   |
| A+ | 2   |
| F  | 2   |

Name: CS-419, dtype: int64

|    |     |
|----|-----|
| A- | 136 |
| A  | 78  |
| B+ | 75  |
| B  | 65  |
| B- | 56  |
| C  | 45  |
| C+ | 39  |
| C- | 29  |
| D+ | 25  |
| D  | 15  |
| F  | 5   |
| A+ | 3   |

Name: CS-423, dtype: int64

|    |     |
|----|-----|
| A- | 236 |
| B+ | 80  |
| A  | 77  |
| B  | 65  |
| B- | 37  |

```

C+      24
C       19
D+       8
C-       7
F        6
D        5
A+       4
W        3
Name: CS-412, dtype: int64
3.019    5
3.058    3
2.793    3
3.443    3
2.206    3
..
2.555    1
2.042    1
2.634    1
2.053    1
1.753    1
Name: CGPA, Length: 491, dtype: int64

```

```

In [18]: #Removing some inconsistent grades
for i in df.columns:
    df.drop(df[(df.loc[:,i]=='WU') | (df.loc[:,i]=='W')].index,inplace=True)
df.reset_index(drop=True,inplace=True)

```

## Observations:

- We have checked the unique values in our dataframe for each column, found some inconsistent grades.
- We have removed those inconsistent grades from the DataFrame.



## Statistical Summary

```
In [19]: df.describe(include="object").T
```

Out[19]:

|                 | count | unique | top      | freq |
|-----------------|-------|--------|----------|------|
| <b>Seat No.</b> | 547   | 547    | CS-97001 | 1    |
| <b>PH-121</b>   | 547   | 12     | A-       | 112  |
| <b>HS-101</b>   | 547   | 12     | A-       | 80   |
| <b>CY-105</b>   | 547   | 11     | A        | 174  |
| <b>HS-105</b>   | 547   | 11     | A        | 95   |
| <b>MT-111</b>   | 547   | 12     | A-       | 105  |
| <b>CS-105</b>   | 547   | 11     | A        | 148  |
| <b>CS-106</b>   | 547   | 11     | A-       | 113  |
| <b>EL-102</b>   | 547   | 11     | A-       | 103  |
| <b>EE-119</b>   | 547   | 11     | A-       | 134  |
| <b>ME-107</b>   | 547   | 11     | A-       | 79   |
| <b>CS-107</b>   | 547   | 11     | A        | 105  |
| <b>HS-205</b>   | 547   | 11     | A-       | 151  |
| <b>MT-222</b>   | 547   | 12     | A-       | 87   |
| <b>EE-222</b>   | 547   | 12     | A-       | 121  |
| <b>MT-224</b>   | 547   | 11     | A-       | 121  |
| <b>CS-210</b>   | 547   | 11     | A-       | 134  |
| <b>CS-211</b>   | 547   | 12     | A-       | 67   |
| <b>CS-203</b>   | 547   | 11     | A-       | 89   |
| <b>CS-214</b>   | 547   | 11     | C        | 75   |
| <b>EE-217</b>   | 547   | 12     | A-       | 138  |
| <b>CS-212</b>   | 547   | 11     | A-       | 102  |
| <b>CS-215</b>   | 547   | 11     | A-       | 78   |
| <b>MT-331</b>   | 547   | 12     | A        | 118  |
| <b>EF-303</b>   | 547   | 11     | B        | 111  |
| <b>HS-304</b>   | 547   | 12     | A-       | 130  |
| <b>CS-301</b>   | 547   | 12     | A-       | 110  |
| <b>CS-302</b>   | 547   | 11     | A-       | 114  |
| <b>TC-383</b>   | 547   | 12     | A        | 106  |
| <b>MT-442</b>   | 547   | 12     | A-       | 141  |
| <b>EL-332</b>   | 547   | 12     | A-       | 96   |
| <b>CS-318</b>   | 547   | 12     | A-       | 89   |
| <b>CS-306</b>   | 547   | 12     | A-       | 120  |
| <b>CS-312</b>   | 547   | 12     | A+       | 94   |
| <b>CS-317</b>   | 547   | 12     | B-       | 81   |
| <b>CS-403</b>   | 547   | 11     | A        | 134  |

|               | count | unique | top | freq |
|---------------|-------|--------|-----|------|
| <b>CS-421</b> | 547   | 12     | B   | 89   |
| <b>CS-406</b> | 547   | 12     | A-  | 252  |
| <b>CS-414</b> | 547   | 12     | A   | 179  |
| <b>CS-419</b> | 547   | 12     | A-  | 123  |
| <b>CS-423</b> | 547   | 12     | A-  | 126  |
| <b>CS-412</b> | 547   | 12     | A-  | 223  |


## Observations:

- "Seat No.": This column contains 547 unique values, indicating that there are 547 distinct seat numbers in the dataset. The most frequent seat number is "CS-97001," which appears once.
- "PH-121": There are 12 unique values in this column. The most common value is "A-" with a frequency of 112.
- "HS-101": Similar to "PH-121," there are 12 unique values, and "A-" is the most frequent value with a frequency of 80.
- "CY-105" and "HS-105": These columns both have 11 unique values, and "A" is the most common value in both, with frequencies of 174 and 95, respectively.
- "MT-111," "CS-105," "CS-106," "EL-102," "EE-119," "ME-107," "CS-107," "HS-205," "MT-222," "EE-222," "MT-224," "CS-210," "CS-211," "CS-203," "CS-214," "EE-217," "CS-212," and "CS-215": Each of these columns has 11 or 12 unique values, and "A-" is the most frequent value in most of them.
- "MT-331" and "HS-304": These columns both have 12 unique values, and "A" is the most common value in both, with frequencies of 118 and 130, respectively.
- "EF-303," "CS-301," and "CS-302": These columns have 11 or 12 unique values, and "A-" is the most common value in each of them.
- "TC-383," "MT-442," "EL-332," "CS-318," "CS-306," "CS-312," "CS-317," "CS-403," "CS-421," "CS-406," "CS-414," "CS-419," and "CS-423": Each of these columns has 11 or 12 unique values, and various grades are the most common values in them.


Overall, it appears that the dataset contains information related to courses, seat numbers, and grades. The grades are categorized with varying frequencies in different columns.

## Visualizing it


```
In [20]: # Bar Chart for grades distribution in each subject
subjects = df.columns
for subject in subjects:
    grades_count = df[subject].value_counts().sort_index()
    plt.bar(grades_count.index, grades_count.values)
    plt.xlabel('Grades')
    plt.ylabel('Count')
    plt.title(f'Grades Distribution in {subject}')
    plt.show()
```




Grades Distribution in PH-121




Grades Distribution in HS-101




Grades Distribution in CY-105




Grades Distribution in HS-105




Grades Distribution in MT-111




Grades Distribution in CS-105



Grades Distribution in CS-106




Grades Distribution in EL-102







Grades Distribution in EE-119




Grades Distribution in ME-107




Grades Distribution in CS-107




Grades Distribution in HS-205




Grades Distribution in MT-222




Grades Distribution in EE-222

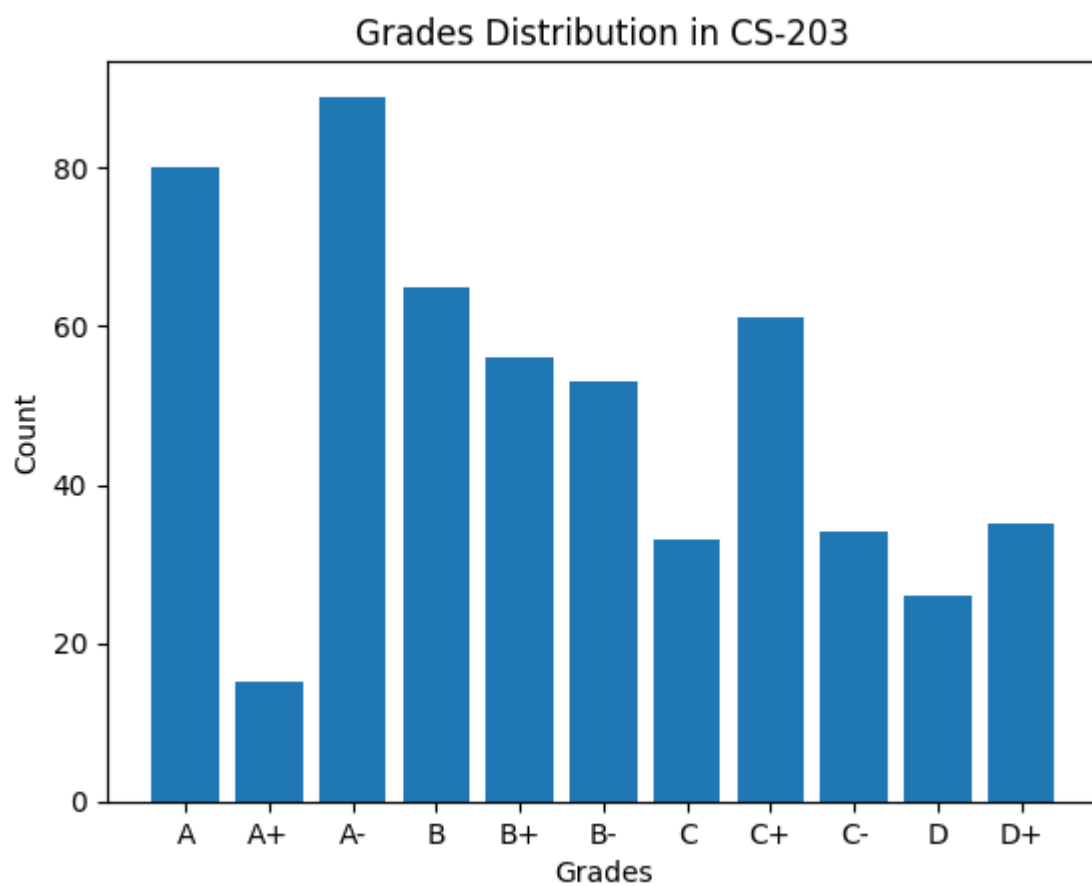
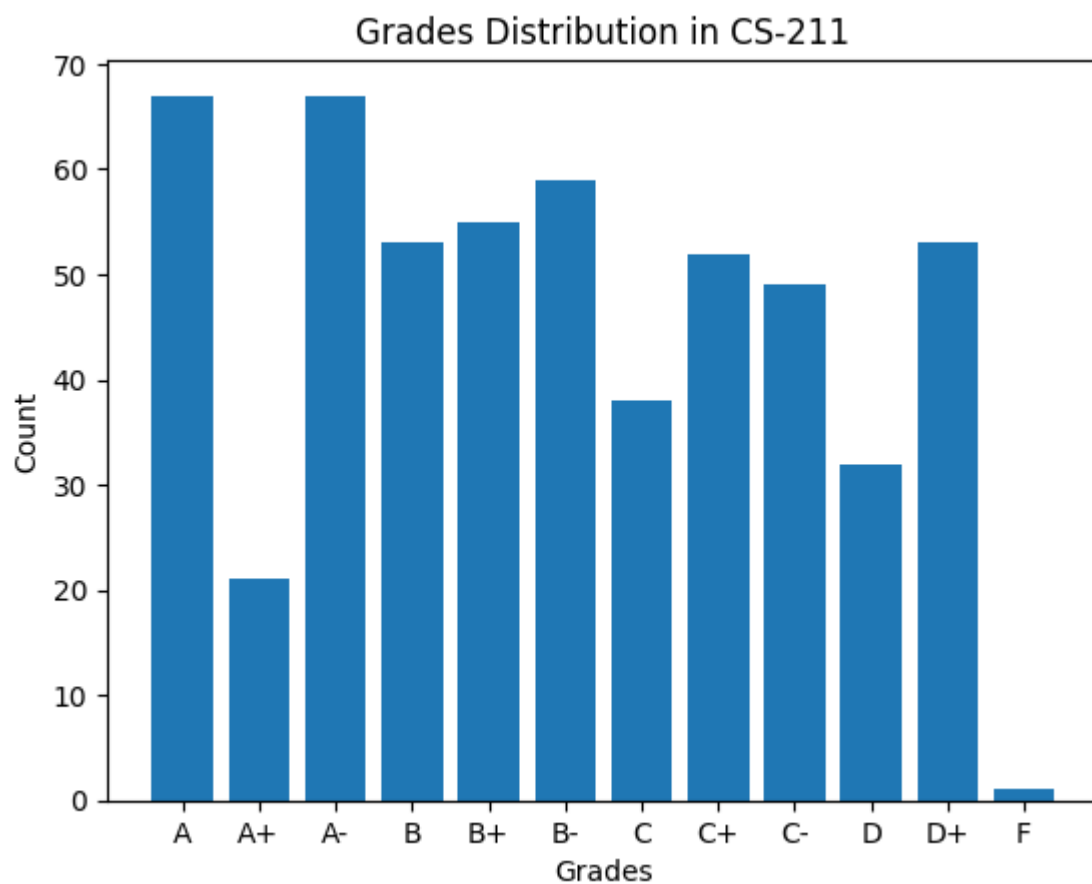


Grades Distribution in MT-224




Grades Distribution in CS-210







Grades Distribution in CS-214




Grades Distribution in EE-217




Grades Distribution in CS-212




Grades Distribution in CS-215



Grades Distribution in MT-331




Grades Distribution in EF-303







Grades Distribution in HS-304




Grades Distribution in CS-301




Grades Distribution in CS-302




Grades Distribution in TC-383




Grades Distribution in MT-442




Grades Distribution in EL-332




Grades Distribution in CS-318




Grades Distribution in CS-306





Grades Distribution in CS-312




Grades Distribution in CS-317







Grades Distribution in CS-406




Grades Distribution in CS-414



Grades Distribution in CS-419




Grades Distribution in CS-423






Grades Distribution in CS-412



Grades Distribution in CGPA



## Dropping Unnecessary column

```
In [21]: df = df.drop('Seat No.', axis=1)
```

## Encoding Categorical Values

```
In [22]: # Define a dictionary to map letter grades to numerical values
grade_mapping = {
    'A+': 4.0,
    'A': 4.0,
    'A-': 3.7,
    'B+': 3.4,
    'B': 3.0,
    'B-': 2.7,
    'C+': 2.4,
    'C': 2.0,
    'C-': 1.7,
    'D+': 1.4,
    'D': 1.0,
    'F': 0.0
}

# Use the replace method to map grades to numerical values for all columns
df = df.replace(grade_mapping)
```

```
In [23]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 547 entries, 0 to 546
Data columns (total 42 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PH-121      547 non-null    float64
1   HS-101      547 non-null    float64
2   CY-105      547 non-null    float64
3   HS-105      547 non-null    float64
4   MT-111      547 non-null    float64
5   CS-105      547 non-null    float64
6   CS-106      547 non-null    float64
7   EL-102      547 non-null    float64
8   EE-119      547 non-null    float64
9   ME-107      547 non-null    float64
10  CS-107      547 non-null    float64
11  HS-205      547 non-null    float64
12  MT-222      547 non-null    float64
13  EE-222      547 non-null    float64
14  MT-224      547 non-null    float64
15  CS-210      547 non-null    float64
16  CS-211      547 non-null    float64
17  CS-203      547 non-null    float64
18  CS-214      547 non-null    float64
19  EE-217      547 non-null    float64
20  CS-212      547 non-null    float64
21  CS-215      547 non-null    float64
22  MT-331      547 non-null    float64
23  EF-303      547 non-null    float64
24  HS-304      547 non-null    float64
25  CS-301      547 non-null    float64
26  CS-302      547 non-null    float64
27  TC-383      547 non-null    float64
28  MT-442      547 non-null    float64
29  EL-332      547 non-null    float64
30  CS-318      547 non-null    float64
31  CS-306      547 non-null    float64
32  CS-312      547 non-null    float64
33  CS-317      547 non-null    float64
34  CS-403      547 non-null    float64
35  CS-421      547 non-null    float64
36  CS-406      547 non-null    float64
37  CS-414      547 non-null    float64
38  CS-419      547 non-null    float64
39  CS-423      547 non-null    float64
40  CS-412      547 non-null    float64
41  CGPA        547 non-null    float64
dtypes: float64(42)
memory usage: 179.6 KB
```

## Observations

- We have dropped "Seat no." from the dataframe as it wont help us in anyway to predict grades
- We have encoded categorical values to numerical values
- Now you can clear see that all columns has float values
- No missing values in our dataframe as count of all the columns are 547
- Memory usage is 179.6KB

```
In [24]: #Statistical summary  
df.describe()
```

```
Out[24]:
```


|              | PH-121     | HS-101     | CY-105     | HS-105     | MT-111     | CS-105     | CS-106     | I     |
|--------------|------------|------------|------------|------------|------------|------------|------------|-------|
| <b>count</b> | 547.000000 | 547.000000 | 547.000000 | 547.000000 | 547.000000 | 547.000000 | 547.000000 | 547.0 |
| <b>mean</b>  | 3.067642   | 2.657952   | 3.310055   | 2.916636   | 3.047166   | 3.366728   | 3.000914   | 3.0   |
| <b>std</b>   | 0.936537   | 0.898506   | 0.873265   | 0.988782   | 0.896047   | 0.765265   | 0.827071   | 0.9   |
| <b>min</b>   | 0.000000   | 0.000000   | 1.000000   | 1.000000   | 0.000000   | 1.000000   | 1.000000   | 1.0   |
| <b>25%</b>   | 2.700000   | 2.000000   | 3.000000   | 2.000000   | 2.400000   | 3.000000   | 2.700000   | 2.4   |
| <b>50%</b>   | 3.400000   | 2.700000   | 3.700000   | 3.000000   | 3.400000   | 3.700000   | 3.000000   | 3.4   |
| <b>75%</b>   | 3.700000   | 3.400000   | 4.000000   | 3.700000   | 3.700000   | 4.000000   | 3.700000   | 3.7   |
| <b>max</b>   | 4.000000   | 4.000000   | 4.000000   | 4.000000   | 4.000000   | 4.000000   | 4.000000   | 4.0   |

8 rows × 42 columns



```
In [25]: # visualization of correlation metrix
plt.figure(figsize=(25,25))
sns.heatmap(df.corr(), annot = True, cmap = 'YlGnBu').set_title('Correlation Fa
```

```
Out[25]: Text(0.5, 1.0, 'Correlation Factors Heat Map')
```



```
In [26]: # correlation with Label
correlation = df.corr()['CGPA'].drop('CGPA')

# Display the correlation values
print(correlation)
```

```
PH-121    0.645401
HS-101    0.556404
CY-105    0.669354
HS-105    0.660944
MT-111    0.738479
CS-105    0.588444
CS-106    0.597217
EL-102    0.651206
EE-119    0.647128
ME-107    0.659035
CS-107    0.601477
HS-205    0.470703
MT-222    0.680591
EE-222    0.577960
MT-224    0.722883
CS-210    0.799214
CS-211    0.672404
CS-203    0.662619
CS-214    0.782109
EE-217    0.699407
CS-212    0.706259
CS-215    0.780897
MT-331    0.550261
EF-303    0.568559
HS-304    0.608500
CS-301    0.755949
CS-302    0.756661
TC-383    0.655092
MT-442    0.626249
EL-332    0.784013
CS-318    0.603264
CS-306    0.758919
CS-312    0.534566
CS-317    0.814143
CS-403    0.639654
CS-421    0.728702
CS-406    0.494329
CS-414    0.370553
CS-419    0.711185
CS-423    0.565454
CS-412    0.495348
Name: CGPA, dtype: float64
```

## Observations:

- We can clearly visualize that all the columns are highly correlated to each other

## Data Splitting

```
In [27]: X = df.drop('CGPA', axis=1) # List of all features  
y = df['CGPA'] # Data of our label
```

```
In [28]: X.shape
```

```
Out[28]: (547, 41)
```

```
In [29]: y.shape
```

```
Out[29]: (547,)
```

### Observations:

- we have splitted the data into features and label.
- Now we have 41 features with 547 observations.
- As we y is our target it has only 574 observations


## Feature Scaling

```
In [30]: scaler = StandardScaler()  
X = pd.DataFrame(scaler.fit_transform(X), columns = X.columns)  
X
```

```
Out[30]:
```

|     | PH-121    | HS-101    | CY-105    | HS-105    | MT-111    | CS-105    | CS-106    | EL-102    | EE-1    |
|-----|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|---------|
| 0   | -0.392914 | -1.401331 | -1.845407 | -0.927884 | -1.504832 | -2.572347 | -2.421491 | -1.420098 | -0.4530 |
| 1   | 0.996450  | -1.846922 | -2.189260 | -1.940155 | -0.387797 | -1.787588 | -2.421491 | 1.054452  | -2.0848 |
| 2   | 0.996450  | 0.381033  | 0.790799  | -0.219295 | 0.394128  | 0.828278  | -0.364164 | 0.408917  | 0.8022  |
| 3   | -2.209774 | -0.287353 | -2.189260 | -1.940155 | -2.286756 | 0.435898  | -1.937414 | -1.420098 | -2.5869 |
| 4   | 0.675827  | 1.160817  | 0.446946  | 0.489295  | 1.064348  | 0.828278  | 0.846029  | 0.408917  | 1.1788  |
| ... | ...       | ...       | ...       | ...       | ...       | ...       | ...       | ...       | ...     |
| 542 | -0.072291 | 1.495011  | 0.790799  | 0.792976  | 1.064348  | 0.828278  | 0.846029  | 0.731685  | 1.1788  |
| 543 | 0.996450  | 1.495011  | 0.790799  | 1.096658  | 1.064348  | 0.828278  | 1.209086  | 0.731685  | 1.1788  |
| 544 | -0.072291 | 1.495011  | 0.446946  | 0.489295  | 1.064348  | 0.828278  | 1.209086  | 1.054452  | 1.1788  |
| 545 | 0.996450  | 0.826624  | -2.647731 | 1.096658  | -2.286756 | -2.572347 | -0.364164 | -1.420098 | -0.4530 |
| 546 | -1.141032 | -1.846922 | -2.647731 | -0.927884 | -1.169721 | -2.572347 | -0.001106 | -0.666974 | -1.3316 |

547 rows × 41 columns



## Checking Best Random State

```
In [31]: maxR2_score=0
maxRS=0
for i in range(1,200):
    X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.33, random_state=i)
    lr=LinearRegression()
    lr.fit(X_train,y_train)
    y_pred=lr.predict(X_test)
    R2=r2_score(y_test,y_pred)
    if R2>maxR2_score:
        maxR2_score=R2
        maxRS=i
print('Best accuracy is', maxR2_score , 'on Random_state', maxRS)
```

Best accuracy is 0.9875788176922384 on Random\_state 46

## Train\_Test\_Split

```
In [32]: x_train,x_test,y_train,y_test= train_test_split(X,y, test_size =0.2, random_state=42)
```



# Model Training & Testing

## LinearRegression

```
In [33]: LR = LinearRegression()

# Perform cross-validation
cv_scores = cross_val_score(LR, x_train, y_train, cv=5, scoring='r2')

# Fit the model on the entire training set
LR.fit(x_train, y_train)

# Predict on the test set
y_pred = LR.predict(x_test)

# Print evaluation metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

# Print cross-validation scores
print('Cross-Validation Scores:', cv_scores)
print('Mean Cross-Validation Score:', cv_scores.mean())
```

R2 Score: 0.9884894119471499  
MAE: 0.05287669778724794  
Cross-Validation Scores: [0.98972109 0.97168207 0.77195307 0.97381632 0.97827864]  
Mean Cross-Validation Score: 0.9370902366556093

## Ridge Regression

```
In [34]: # Define the parameter grid
param_grid = {'alpha': [0.1, 1, 10, 100]} # Example parameter values, you can

# Create the Ridge regression model
R = Ridge()

# Perform grid search with cross-validation
grid_search = GridSearchCV(R, param_grid, cv=5, scoring='r2')
grid_search.fit(x_train, y_train)

# Print the best parameters
print('Best Parameters:', grid_search.best_params_)

# Fit the model with the best parameters on the entire training set
best_R = grid_search.best_estimator_
best_R.fit(x_train, y_train)

# Predict on the test set
y_pred = best_R.predict(x_test)

# Print evaluation metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

# Perform cross-validation with the best model
cv_scores = cross_val_score(best_R, x_train, y_train, cv=5, scoring='r2')

# Print cross-validation scores
print('Cross-Validation Scores:', cv_scores)
print('Mean Cross-Validation Score:', cv_scores.mean())
```

Best Parameters: {'alpha': 100}

R2 Score: 0.992129009223353

MAE: 0.04087843834677456

Cross-Validation Scores: [0.99406179 0.98239821 0.78052778 0.98514285 0.99237896]

Mean Cross-Validation Score: 0.9469019170790579

## Lasso Regression

```
In [35]: # Define the parameter grid
param_grid = {'alpha': [0.1, 1, 10, 100]} # Example parameter values, you can

# Create the Lasso regression model
L = Lasso()

# Perform grid search with cross-validation
grid_search = GridSearchCV(L, param_grid, cv=5, scoring='r2')
grid_search.fit(x_train, y_train)

# Print the best parameters
print('Best Parameters:', grid_search.best_params_)

# Fit the model with the best parameters on the entire training set
best_L = grid_search.best_estimator_
best_L.fit(x_train, y_train)

# Predict on the test set
y_pred = best_L.predict(x_test)

# Print evaluation metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

# Perform cross-validation with the best model
cv_scores = cross_val_score(best_L, x_train, y_train, cv=5, scoring='r2')

# Print cross-validation scores
print('Cross-Validation Scores:', cv_scores)
print('Mean Cross-Validation Score:', cv_scores.mean())
```

Best Parameters: {'alpha': 0.1}

R2 Score: 0.8979205182363923

MAE: 0.16203976784773086

Cross-Validation Scores: [0.90262411 0.88373868 0.72248202 0.87741962 0.90261066]

Mean Cross-Validation Score: 0.8577750169235688

## Decision Tree Regressor

```
In [36]: # Define the parameter grid
param_grid = {
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': [None, 'sqrt', 'log2']
}

# Create the Decision Tree Regressor model
DTR = DecisionTreeRegressor()

# Perform grid search with cross-validation
grid_search = GridSearchCV(DTR, param_grid, cv=5, scoring='r2')
grid_search.fit(x_train, y_train)

# Print the best parameters
print('Best Parameters:', grid_search.best_params_)

# Fit the model with the best parameters on the entire training set
best_DTR = grid_search.best_estimator_
best_DTR.fit(x_train, y_train)

# Predict on the test set
y_pred = best_DTR.predict(x_test)

# Print evaluation metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

# Perform cross-validation with the best model
cv_scores = cross_val_score(best_DTR, x_train, y_train, cv=5, scoring='r2')

# Print cross-validation scores
print('Cross-Validation Scores:', cv_scores)
print('Mean Cross-Validation Score:', cv_scores.mean())

Best Parameters: {'max_depth': 10, 'max_features': 'sqrt', 'min_samples_leaf': 2, 'min_samples_split': 2}
R2 Score: 0.8041219851751172
MAE: 0.18967030893349074
Cross-Validation Scores: [0.84960142 0.82536029 0.62495131 0.8751598 0.80308116]
Mean Cross-Validation Score: 0.7956307949264849
```

## Random Forest Regressor

```
In [37]: # Define the parameter grid
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt']
}

# Create the Random Forest Regressor model
RFR = RandomForestRegressor()

# Perform grid search with cross-validation
grid_search = GridSearchCV(RFR, param_grid, cv=5, scoring='r2')
grid_search.fit(x_train, y_train)

# Print the best parameters
print('Best Parameters:', grid_search.best_params_)

# Fit the model with the best parameters on the entire training set
best_RFR = grid_search.best_estimator_
best_RFR.fit(x_train, y_train)

# Predict on the test set
y_pred = best_RFR.predict(x_test)

# Print evaluation metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

# Perform cross-validation with the best model
cv_scores = cross_val_score(best_RFR, x_train, y_train, cv=5, scoring='r2')

# Print cross-validation scores
print('Cross-Validation Scores:', cv_scores)
print('Mean Cross-Validation Score:', cv_scores.mean())

Best Parameters: {'max_depth': 10, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 200}
R2 Score: 0.9430268670773639
MAE: 0.09561351183951926
Cross-Validation Scores: [0.96510473 0.95120716 0.73974802 0.96402785 0.97202262]
Mean Cross-Validation Score: 0.9184220773387425
```

## Extra Tree Regressor

```
In [38]: # Define the parameter grid
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt']
}

# Create the Extra Trees Regressor model
ETR = ExtraTreesRegressor()

# Perform grid search with cross-validation
grid_search = GridSearchCV(ETR, param_grid, cv=5, scoring='r2')
grid_search.fit(x_train, y_train)

# Print the best parameters
print('Best Parameters:', grid_search.best_params_)

# Fit the model with the best parameters on the entire training set
best_ETR = grid_search.best_estimator_
best_ETR.fit(x_train, y_train)

# Predict on the test set
y_pred = best_ETR.predict(x_test)

# Print evaluation metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

# Perform cross-validation with the best model
cv_scores = cross_val_score(best_ETR, x_train, y_train, cv=5, scoring='r2')

# Print cross-validation scores
print('Cross-Validation Scores:', cv_scores)
print('Mean Cross-Validation Score:', cv_scores.mean())

Best Parameters: {'max_depth': 20, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 300}
R2 Score: 0.9545728090941457
MAE: 0.08417658879268887
Cross-Validation Scores: [0.97292878 0.95777007 0.74884279 0.9704137 0.97293474]
Mean Cross-Validation Score: 0.9245780143030384
```

## Ada Boost Regressor

```
In [39]: # Define the parameter grid
param_grid = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.1, 0.01, 0.001],
    'max_depth': [3, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt']
}

# Create the Gradient Boosting Regressor model
GBR = GradientBoostingRegressor()

# Perform grid search with cross-validation
grid_search = GridSearchCV(GBR, param_grid, cv=5, scoring='r2')
grid_search.fit(x_train, y_train)

# Print the best parameters
print('Best Parameters:', grid_search.best_params_)

# Fit the model with the best parameters on the entire training set
best_GBR = grid_search.best_estimator_
best_GBR.fit(x_train, y_train)

# Predict on the test set
y_pred = best_GBR.predict(x_test)

# Print evaluation metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

# Perform cross-validation with the best model
cv_scores = cross_val_score(best_GBR, x_train, y_train, cv=5, scoring='r2')

# Print cross-validation scores
print('Cross-Validation Scores:', cv_scores)
print('Mean Cross-Validation Score:', cv_scores.mean())

Best Parameters: {'learning_rate': 0.1, 'max_depth': 3, 'max_features': 'sqrt', 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 200}
R2 Score: 0.9624314018087885
MAE: 0.08048523532908322
Cross-Validation Scores: [0.97537463 0.95857745 0.77329232 0.97042696 0.96666871]
Mean Cross-Validation Score: 0.9288680150630805
```

## XGB Regressor

```
In [41]: # Define the parameter grid
param_grid = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.1, 0.01, 0.001],
    'max_depth': [3, 5, 10],
    'min_child_weight': [1, 3, 5],
    'gamma': [0, 0.1, 0.2],
    'subsample': [0.8, 1.0],
    'colsample_bytree': [0.8, 1.0]
}

# Create the XGBoost Regressor model
XGBR = XGBRegressor()

# Perform grid search with cross-validation
grid_search = GridSearchCV(XGBR, param_grid, cv=5, scoring='r2')
grid_search.fit(x_train, y_train)

# Print the best parameters
print('Best Parameters:', grid_search.best_params_)

# Fit the model with the best parameters on the entire training set
best_XGBR = grid_search.best_estimator_
best_XGBR.fit(x_train, y_train)

# Predict on the test set
y_pred = best_XGBR.predict(x_test)

# Print evaluation metrics
print('R2 Score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

# Perform cross-validation with the best model
cv_scores = cross_val_score(best_XGBR, x_train, y_train, cv=5, scoring='r2')

# Print cross-validation scores
print('Cross-Validation Scores:', cv_scores)
print('Mean Cross-Validation Score:', cv_scores.mean())

Best Parameters: {'colsample_bytree': 0.8, 'gamma': 0, 'learning_rate': 0.1,
 'max_depth': 3, 'min_child_weight': 1, 'n_estimators': 300, 'subsample': 0.8}
R2 Score: 0.9680914803790569
MAE: 0.06646585891897029
Cross-Validation Scores: [0.97257635 0.947726  0.75845435 0.96550984 0.96851
 745]
Mean Cross-Validation Score: 0.9225568005070667
```



## Saving The Model

```
In [42]: filename = 'LinearRegression.pkl'
pickle.dump(LR, open(filename, 'wb'))
pickle.dump(scaler, open('scaler.pkl', 'wb'))
```

```
In [43]: with open('LinearRegression.pkl', 'rb') as file:
        data = pickle.load(file)
```

```
In [44]: data.predict(X_test)
```

```
Out[44]: array([3.77068565, 3.75407909, 2.54638351, 2.09360857, 3.79929245,
                2.94961198, 2.09316688, 3.23065178, 2.9072488 , 2.03583939,
                3.62151759, 3.01708529, 3.54714463, 2.80159418, 3.42954448,
                3.7285125 , 3.14013385, 3.94642353, 2.82604138, 2.07526458,
                2.63434344, 3.08805893, 2.81268765, 3.55520011, 3.48156609,
                3.23621862, 3.20573641, 1.97585951, 3.1965465 , 3.19690244,
                3.32649891, 3.08564728, 2.35530664, 2.5222813 , 3.37919198,
                3.70367376, 3.51162938, 2.74767634, 1.89004226, 2.75939195,
                3.3905639 , 2.91442493, 2.2607173 , 2.79347019, 3.59859323,
                2.64477587, 2.11762028, 2.91221773, 3.36521106, 3.34981952,
                3.70027887, 3.21621348, 3.71687807, 3.43379881, 2.05197933,
                2.70001763, 2.80321349, 3.68061392, 2.75416651, 3.44985303,
                2.01609771, 2.49585468, 3.22180723, 2.66438405, 3.28641425,
                2.69450312, 3.18749046, 2.69617306, 3.7197727 , 2.79241047,
                2.68502294, 2.6333894 , 2.31021338, 3.85124306, 2.21982604,
                1.67976166, 3.22778617, 3.42832287, 2.60992268, 3.5804398 ,
                3.91321072, 2.36095309, 3.23243762, 3.64396685, 2.37045648,
                2.95291942, 3.01967734, 3.40840739, 3.73934669, 3.51014454,
                3.45246303, 3.33509822, 3.31931864, 2.88377436, 3.02355165,
                2.39303435, 2.34369662, 3.65112597, 2.07028116, 3.37433738,
                3.34307375, 2.97092875, 3.23823953, 2.95603269, 2.25407779,
                2.25122401, 2.71968032, 3.42939103, 2.81673461, 3.52555136,
                3.28047123, 2.4037378 , 3.33179508, 2.50166223, 2.44402452,
                3.4338544 , 2.97940816, 3.28026008, 3.72369396, 3.75522856,
                2.60246276, 2.79758156, 3.06762496, 3.07697157, 2.65655002,
                2.66119539, 2.510536 , 2.64403298, 3.3414724 , 3.31516961,
                2.26097537, 1.94256418, 2.11220247, 2.09339029, 2.50179947,
                2.78796 , 3.08631393, 1.88918392, 1.77854538, 3.52114735,
                2.92457503, 2.93211482, 3.15352026, 3.00933502, 2.49063925,
                3.3524095 , 2.37224261, 2.41506061, 3.04507289, 3.66893984,
                3.17324682, 3.21287961, 3.6440449 , 3.56021996, 2.89916625,
                2.75021942, 3.43468634, 2.58246441, 3.87106498, 3.576756 ,
                3.68477213, 2.98628938, 2.68814171, 3.46140601, 2.60150606,
                3.28057787, 2.72198083, 2.09389936, 3.79780698, 3.90488737,
                1.95304486, 3.10633644, 2.32663802, 3.01475513, 2.43784521,
                2.10324206, 2.37956128, 2.8256216 , 2.91374611, 3.63339214,
                3.26468148])
```

```
In [ ]:
```

